

EEET2574 | Big Data for Engineering

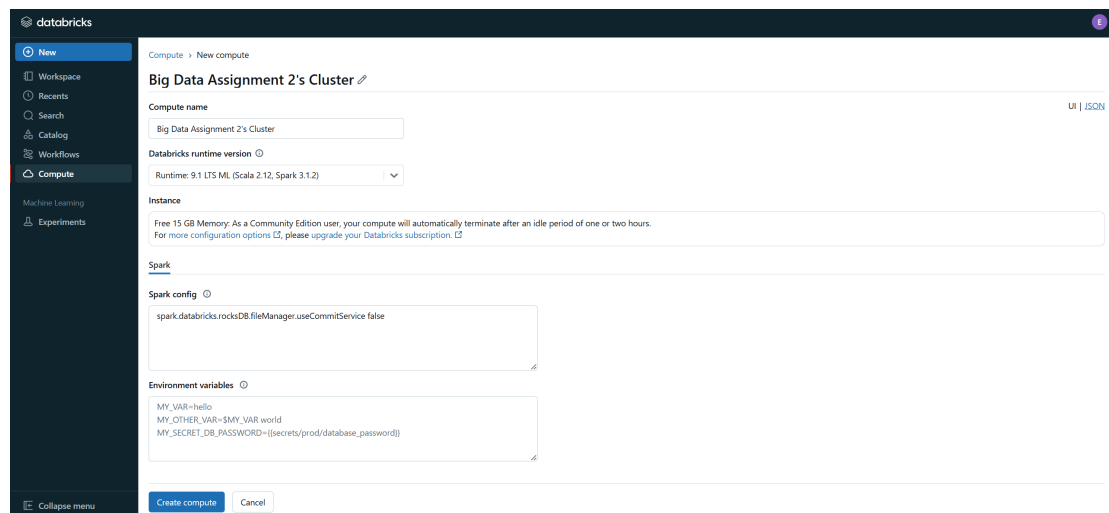
Assignment 2: MongoDB and Spark

- **Author:** Tran Manh Cuong - s3974735
- **Instructor:** Dr. Arthur Tang

Task 0: Instruction on how to run the notebook

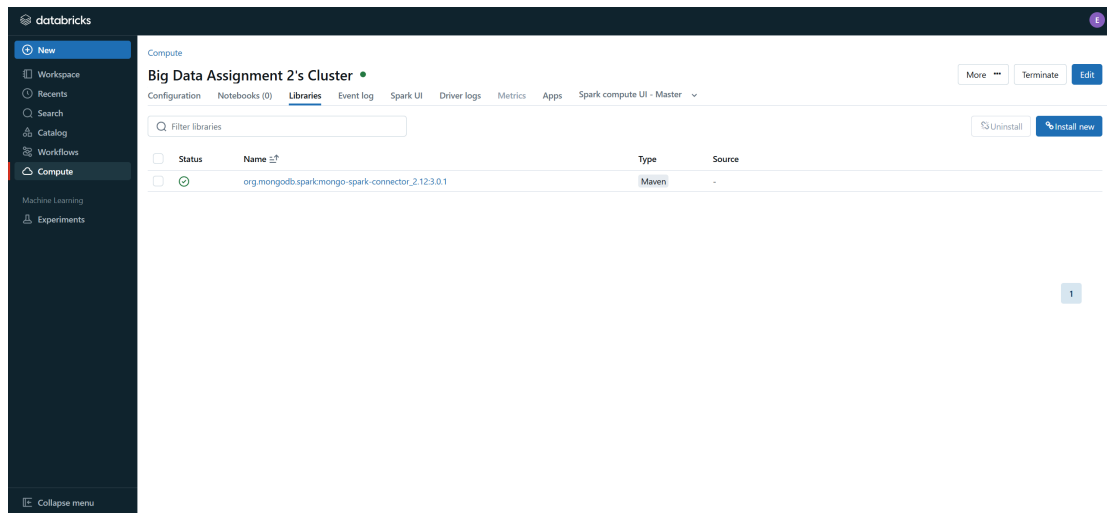
Step 1: Create a compute cluster:

- On the sidebar, select **Compute**.
- Select **Create compute**.
- Enter compute name: **Big Data Assignment 2's Cluster**.
- Choose the databricks runtime version: **9.1 LTS ML (Scala 2.12, Spark 3.1.2)**.
- Select **Create compute**.



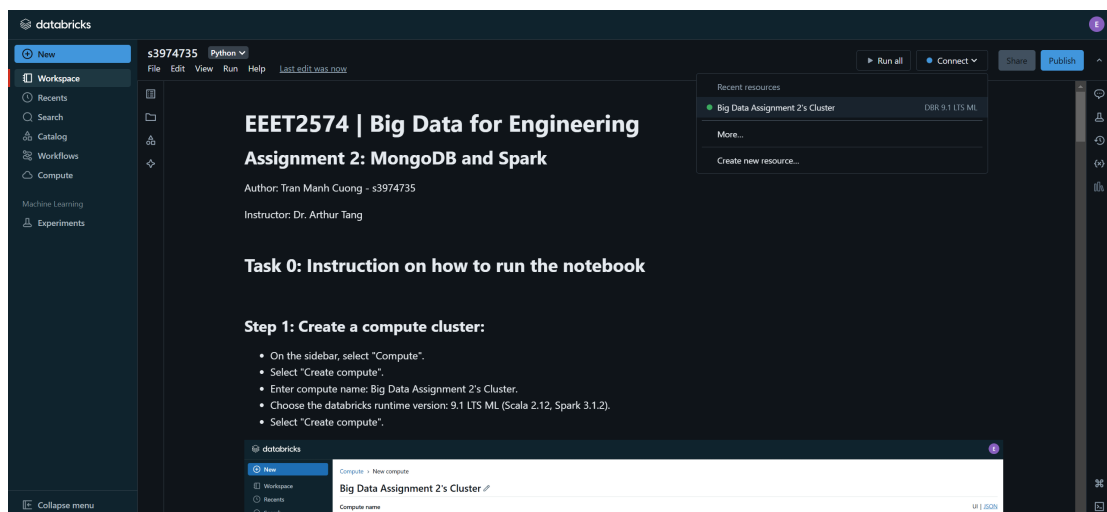
Step 2: Install mongodb spark connector library

- On the navigation bar, select **Libraries**.
- Select **Install new**.
- For Library Source, select **Maven**.
- For Coordinates, select **Select Packages**.
- Select **Spark Packages**.
- Search and select **mongo-spark** with version **3.0.1**.
- Select **Install**.



Step 3: Attach cluster to notebook

- Import the notebook to databricks using .dbc or .ipynb file.
- Select the notebook.
- For the Connect, select the created cluster.



Task 1: MongoDB

Load Datasets to MongoDB Collections

The datasets had already been loaded into MongoDB, so you don't need to run this code.

```
In [0]: import pymongo
import pandas as pd
import os

# You can replace the data directory with your own path
ELECTRICITY_DIR = "./data/Electricity"
GAS_DIR = "./data/Gas"
```

```

# You can replace the URI with your own MongoDB URI
MONGO_URI = "mongodb+srv://cuongtran:cuongtran123@cluster0.5zjcy.mongodb.net/?re

# MongoDB connection
client = pymongo.MongoClient(MONGO_URI)
db = client["main"]

# Data mapping for collections
collections_mapping = {
    "electricity": {
        "2018": ["coteq_electricity_2018.csv", "stedin_electricity_2018.csv", "w
        "2019": ["coteq_electricity_2019.csv", "stedin_electricity_2019.csv", "w
        "2020": ["coteq_electricity_2020.csv", "stedin_electricity_2020.csv", "w
    },
    "gas": {
        "2018": ["coteq_gas_2018.csv", "stedin_gas_2018.csv", "westland-infra_ga
        "2019": ["coteq_gas_2019.csv", "stedin_gas_2019.csv", "westland-infra_ga
        "2020": ["coteq_gas_2020.csv", "stedin_gas_2020.csv", "westland-infra_ga
    }
}

# Check if a collection exists in MongoDB
def is_collection_exist(collection_name):
    return collection_name in db.list_collection_names()

# Load data into MongoDB
def load_data_to_mongodb(database, directory, year, file_list):
    collection_name = f"{database}_{year}"
    if is_collection_exist(collection_name):
        print(f"Collection {collection_name} already exists. Skipping data load
        return

    collection = db[collection_name]
    for file_name in file_list:
        file_path = os.path.join(directory, file_name)
        df = pd.read_csv(file_path)
        records = df.to_dict(orient="records")
        collection.insert_many(records)
        print(f"Data from {file_name} has been loaded to collection {collection_

# Main function
def main():
    for database, years in collections_mapping.items():
        for year, file_list in years.items():
            if database == "electricity":
                load_data_to_mongodb(database, ELECTRICITY_DIR, year, file_list)
            elif database == "gas":
                load_data_to_mongodb(database, GAS_DIR, year, file_list)

if __name__ == "__main__":
    main()

```

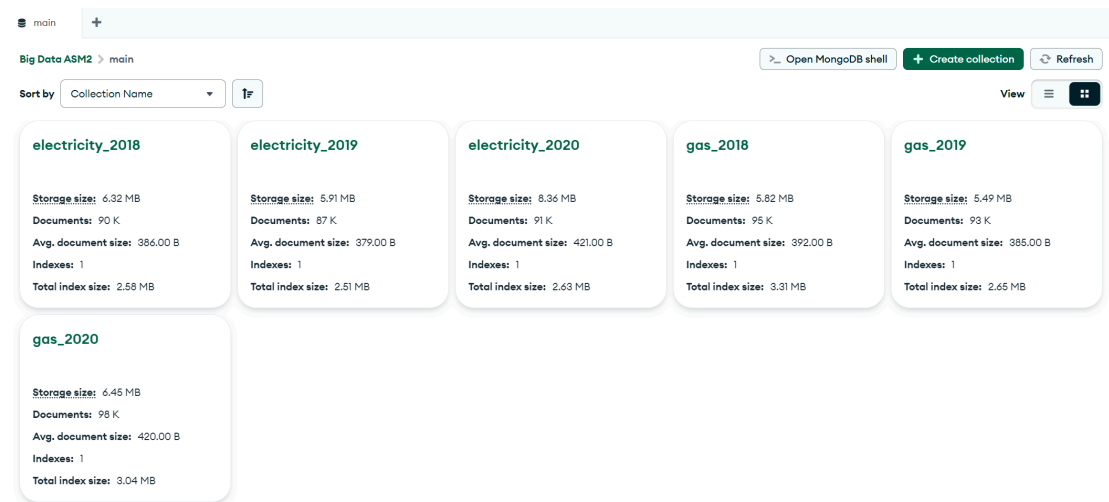
Question 1A: How many collections do you have? Why?

The MongoDB database contains a total of 6 collections, categorized based on type of data (Electricity, Gas) and year (2018, 2019, 2020):

- Electricity collections: electricity_2018, electricity_2019, electricity_2020
- Gas collections: gas_2018, gas_2019, gas_2020

This design choice allows targeted querying and manipulation during data ingestion and cleaning in Task 2. By having distinct collections for each category and year, we can streamline the transformation pipeline, as each collection can be processed individually or combined as needed. Furthermore, this structure supports scalable data management. If new data for additional years or companies becomes available, it can easily be incorporated into separate collections without affecting the existing ones.

Another reason for this design is to simplify model training and tracking in Task 3. Since the task requires using 2018 and 2019 data for training and 2020 for testing, the separate collections make it straightforward to load and partition data for machine learning pipelines. Additionally, visualizations in Task 4 benefit from this schema as it allows flexibility to aggregate data across years or focus on specific time frames and data types.



Task 2: Data ingestion and data cleaning/transformation

Data Ingestion

Load electricity and gas consumption data from MongoDB into PySpark DataFrames

```
In [0]: from pyspark.sql import SparkSession

# MongoDB URI
MONGO_URI = "mongodb+srv://cuongtran:cuongtran123@cluster0.5zjcy.mongodb.net/?re

# Initialize Spark session
```

```

spark = SparkSession.builder \
    .appName("Big_Data_ASM2") \
    .config("spark.mongodb.input.uri", MONGO_URI) \
    .config("spark.mongodb.output.uri", MONGO_URI) \
    .getOrCreate()

# Electricity collections
electricity_2018_df = spark.read.format("com.mongodb.spark.sql.DefaultSource") \
    .option("uri", MONGO_URI) \
    .option("database", "main") \
    .option("collection", "electricity_2018") \
    .load()

electricity_2019_df = spark.read.format("com.mongodb.spark.sql.DefaultSource") \
    .option("uri", MONGO_URI) \
    .option("database", "main") \
    .option("collection", "electricity_2019") \
    .load()

electricity_2020_df = spark.read.format("com.mongodb.spark.sql.DefaultSource") \
    .option("uri", MONGO_URI) \
    .option("database", "main") \
    .option("collection", "electricity_2020") \
    .load()

# Gas collections
gas_2018_df = spark.read.format("com.mongodb.spark.sql.DefaultSource") \
    .option("uri", MONGO_URI) \
    .option("database", "main") \
    .option("collection", "gas_2018") \
    .load()

gas_2019_df = spark.read.format("com.mongodb.spark.sql.DefaultSource") \
    .option("uri", MONGO_URI) \
    .option("database", "main") \
    .option("collection", "gas_2019") \
    .load()

gas_2020_df = spark.read.format("com.mongodb.spark.sql.DefaultSource") \
    .option("uri", MONGO_URI) \
    .option("database", "main") \
    .option("collection", "gas_2020") \
    .load()

```

Concatenate the 2018 and 2019 dataframes to create train datasets and assigns the 2020 dataframes as test datasets for Task 3.

```

In [0]: # Train datasets
electricity_train_df = electricity_2018_df.union(electricity_2019_df)
gas_train_df = gas_2018_df.union(gas_2019_df)

# Test datasets
electricity_test_df = electricity_2020_df
gas_test_df = gas_2020_df

```

Data Exploration

Convert the Spark DataFrames for electricity and gas train datasets into Pandas DataFrames using the `toPandas()` method. The purpose of this conversion is to facilitate data exploration using Pandas' versatile and intuitive tools, which are better suited for detailed analysis, such as summarizing data, checking for missing values, and visualizing distributions

```
In [0]: electricity_pandas_df = electricity_train_df.toPandas()
gas_pandas_df = gas_train_df.toPandas()
```

Display the number of rows and columns to understand the overall size of the dataset.

```
In [0]: electricity_pandas_df.shape
```

```
Out[69]: (176807, 16)
```

```
In [0]: gas_pandas_df.shape
```

```
Out[70]: (187606, 16)
```

Preview the first five rows of the electricity and gas train datasets to understand the structure, column names, and sample data.

```
In [0]: electricity_pandas_df.head()
```

	%Defintieve aansl (NRM)	_id	annual_consume	annual_consume_lowtarif_p
0	NaN	{'oid': '676d0e00999375116ce42b08'}	4122	89
1	NaN	{'oid': '676d0e00999375116ce42b09'}	1800	94
2	NaN	{'oid': '676d0e00999375116ce42b0a'}	1315	100
3	NaN	{'oid': '676d0e00999375116ce42b0b'}	6379	92
4	NaN	{'oid': '676d0e00999375116ce42b0c'}	4404	92

```
In [0]: gas_pandas_df.head()
```

	%Defintieve aansl (NRM)	_id	annual_consume	annual_consume_lowtarif_p
0	NaN	{'oid': '676d0e57999375116ce842fd'}	3457	
1	NaN	{'oid': '676d0e57999375116ce842fe'}	4036	
2	NaN	{'oid': '676d0e57999375116ce842ff'}	3695	
3	NaN	{'oid': '676d0e57999375116ce84300'}	3307	
4	NaN	{'oid': '676d0e57999375116ce84301'}	2306	

Display the data types of each column to understand which features are numerical, categorical, or require further transformation. The result indicates that most numerical columns are suitable for scaling, while categorical columns will need encoding during the data transformation process.

```
In [0]: electricity_pandas_df.dtypes
```

```
Out[12]: %Defintieve aansl (NRM)          float64
_id                                         object
annual_consume                           int32
annual_consume_lowtarif_perc             float64
city                                       object
delivery_perc                           float64
net_manager                              object
num_connections                          int32
perc_of_active_connections               float64
purchase_area                            object
smartmeter_perc                         float64
street                                    object
type_conn_perc                          float64
type_of_connection                       object
zipcode_from                             object
zipcode_to                               object
dtype: object
```

```
In [0]: gas_pandas_df.dtypes
```

```

Out[13]: %Defintieve aansl (NRM)      float64
_id                                     object
annual_consume                        int32
annual_consume_lowtarif_perc         float64
city                                  object
delivery_perc                        float64
net_manager                          object
num_connections                      int32
perc_of_active_connections            float64
purchase_area                        object
smartmeter_perc                     float64
street                               object
type_conn_perc                      float64
type_of_connection                   object
zipcode_from                         object
zipcode_to                           object
dtype: object

```

Check for missing values in each column to identify potential data quality issues. The results reveal that the **%Defintieve aansl (NRM)** column has a significant number of missing values (171,965 in electricity and 182,994 in gas datasets), indicating it may require removal or imputation. All other columns have no missing values, suggesting they are ready for further analysis and transformation.

```
In [0]: electricity_pandas_df.isnull().sum()
```

```

Out[47]: %Defintieve aansl (NRM)      171965
_id                                     0
annual_consume                        0
annual_consume_lowtarif_perc         0
city                                  0
delivery_perc                        0
net_manager                          0
num_connections                      0
perc_of_active_connections            0
purchase_area                        0
smartmeter_perc                     0
street                               0
type_conn_perc                      0
type_of_connection                   0
zipcode_from                         0
zipcode_to                           0
dtype: int64

```

```
In [0]: gas_pandas_df.isnull().sum()
```



```

Out[15]: %Defintieve aansl (NRM)      182994
_id                                     0
annual_consume                        0
annual_consume_lowtarif_perc         0
city                                  0
delivery_perc                        0
net_manager                          0
num_connections                      0
perc_of_active_connections           0
purchase_area                       0
smartmeter_perc                     0
street                              0
type_conn_perc                      0
type_of_connection                   0
zipcode_from                        0
zipcode_to                          0
dtype: int64

```

Identify and categorize columns into numerical and categorical features for targeted exploration and transformation.

```

In [0]: numerical_columns = [
        "annual_consume",
        "annual_consume_lowtarif_perc",
        "delivery_perc",
        "num_connections",
        "perc_of_active_connections",
        "smartmeter_perc",
        "type_conn_perc"
    ]

    categorical_columns = [
        "city",
        "net_manager",
        "purchase_area",
        "street",
        "type_of_connection",
        "zipcode_from",
        "zipcode_to"
    ]

```

Generate summary statistics for numerical columns to understand their distribution, central tendency, and spread. The electricity dataset shows notable outliers in **annual_consume** (max 110,857 vs. mean 3,913) and **num_connections** (max 1,146 vs. mean 25), indicating skewness that may require log transformation. Similarly, the gas dataset reveals significant skewness in **annual_consume** (max 27,917 vs. mean 1,559) and **num_connections** (max 1,065 vs. mean 22). These findings highlight the need for outlier handling, scaling, and possible transformation during the data cleaning and transformation phases. Additionally, the constant value in **annual_consume_lowtarif_perc** for gas indicates it may not provide significant variability for modeling.

```

In [0]: electricity_pandas_df[numerical_columns].describe()

```

	annual_consume	annual_consume_lowtarif_perc	delivery_perc	num_connections
count	176807.000000	176807.000000	176807.000000	176807.000000
mean	3913.422755	87.088556	94.702266	25.088967
std	3281.350003	18.725204	10.108968	18.562825
min	0.000000	0.000000	0.000000	10.000000
25%	2387.000000	81.480000	92.860000	16.000000
50%	3114.000000	95.240000	100.000000	21.000000
75%	4055.000000	100.000000	100.000000	28.000000
max	110857.000000	100.000000	100.000000	1146.000000

In [0]: `gas_pandas_df[numerical_columns].describe()`

	annual_consume	annual_consume_lowtarif_perc	delivery_perc	num_connections
count	187606.000000	187606.0	187606.000000	187606.000000
mean	1559.704295	0.0	99.999916	22.961163
std	959.909904	0.0	0.022761	14.325173
min	0.000000	0.0	90.910000	10.000000
25%	1065.000000	0.0	100.000000	15.000000
50%	1374.000000	0.0	100.000000	20.000000
75%	1825.000000	0.0	100.000000	26.000000
max	27917.000000	0.0	100.000000	1065.000000

Analyze the unique value counts for categorical columns to understand their diversity and identify potential encoding challenges. Both datasets have a high cardinality in **street**, **zipcode_from**, and **zipcode_to**, which may require dimensionality reduction or exclusion if they do not provide meaningful insights. Columns like **net_manager**, **purchase_area**, and **type_of_connection** have relatively fewer unique values, making them suitable for encoding during the transformation process.

In [0]: `electricity_pandas_df[categorical_columns].nunique()`

```
Out[19]: city                287
net_manager                10
purchase_area              9
street                 29426
type_of_connection        10
zipcode_from             89513
zipcode_to               88945
dtype: int64
```

In [0]: `gas_pandas_df[categorical_columns].nunique()`

```

Out[20]: city                491
net_manager                14
purchase_area              39
street                   31598
type_of_connection         6
zipcode_from             95617
zipcode_to                95087
dtype: int64

```

Analyze the distribution of numerical columns in both datasets using histograms to identify data spread, skewness, and potential outliers. The visualizations reveal that columns like **annual_consume**, **num_connections**, and **type_conn_perc** exhibit highly skewed distributions, indicating the need for transformation techniques like log scaling to reduce skewness. Additionally, **smartmeter_perc** and **delivery_perc** show a more uniform distribution, which can be directly utilized in model training without transformations.

```

In [0]: import matplotlib.pyplot as plt


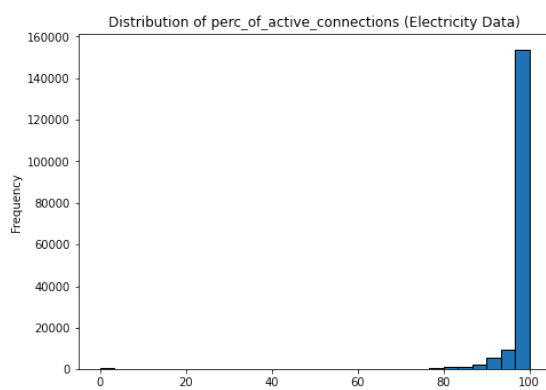
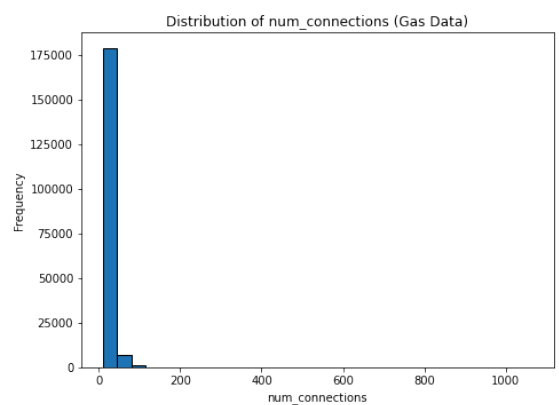
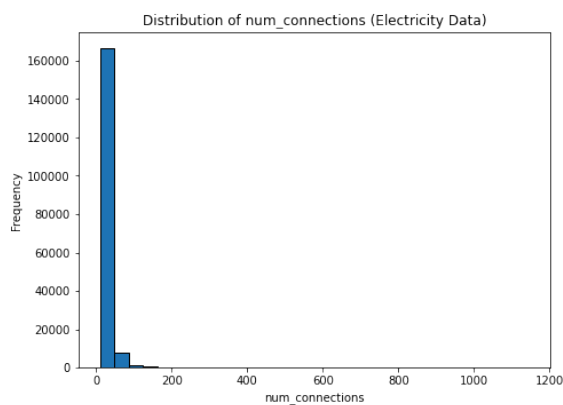
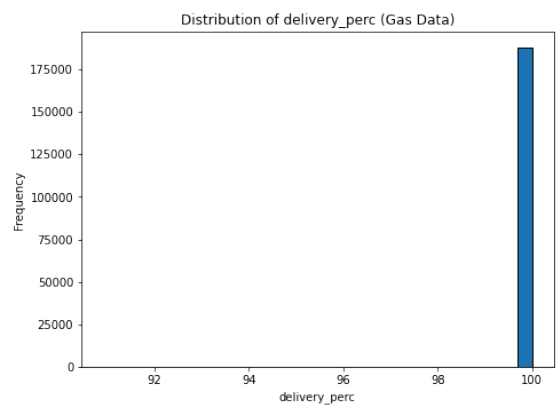
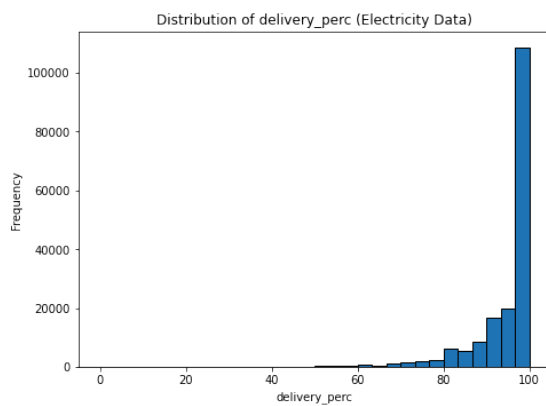
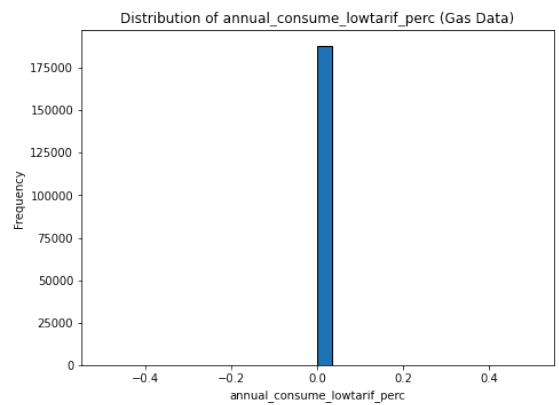
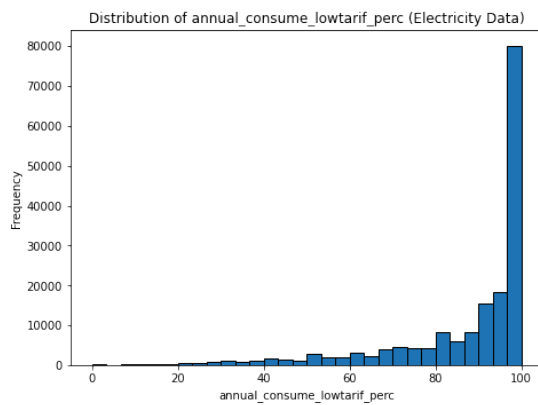
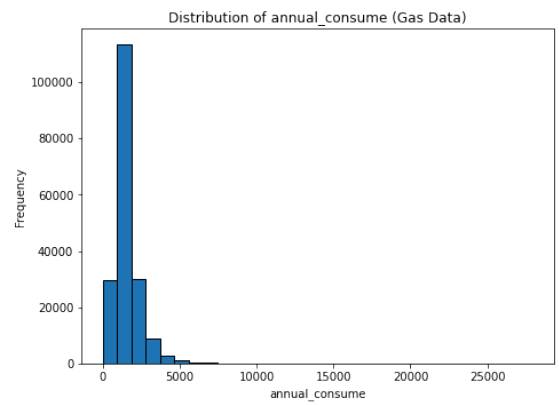
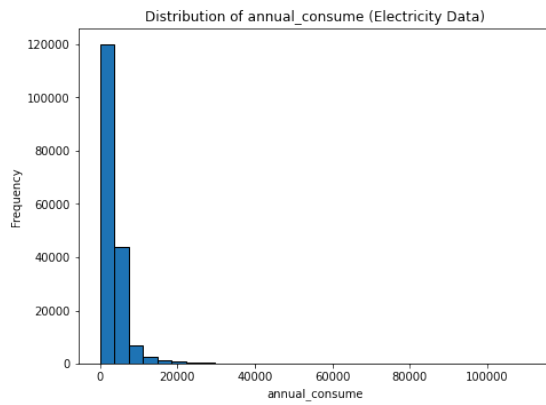
fig, axes = plt.subplots(len(numerical_columns), 2, figsize=(14, 5 * len(numerical_columns)))
fig.tight_layout(w_pad=10.0, h_pad=5.0)

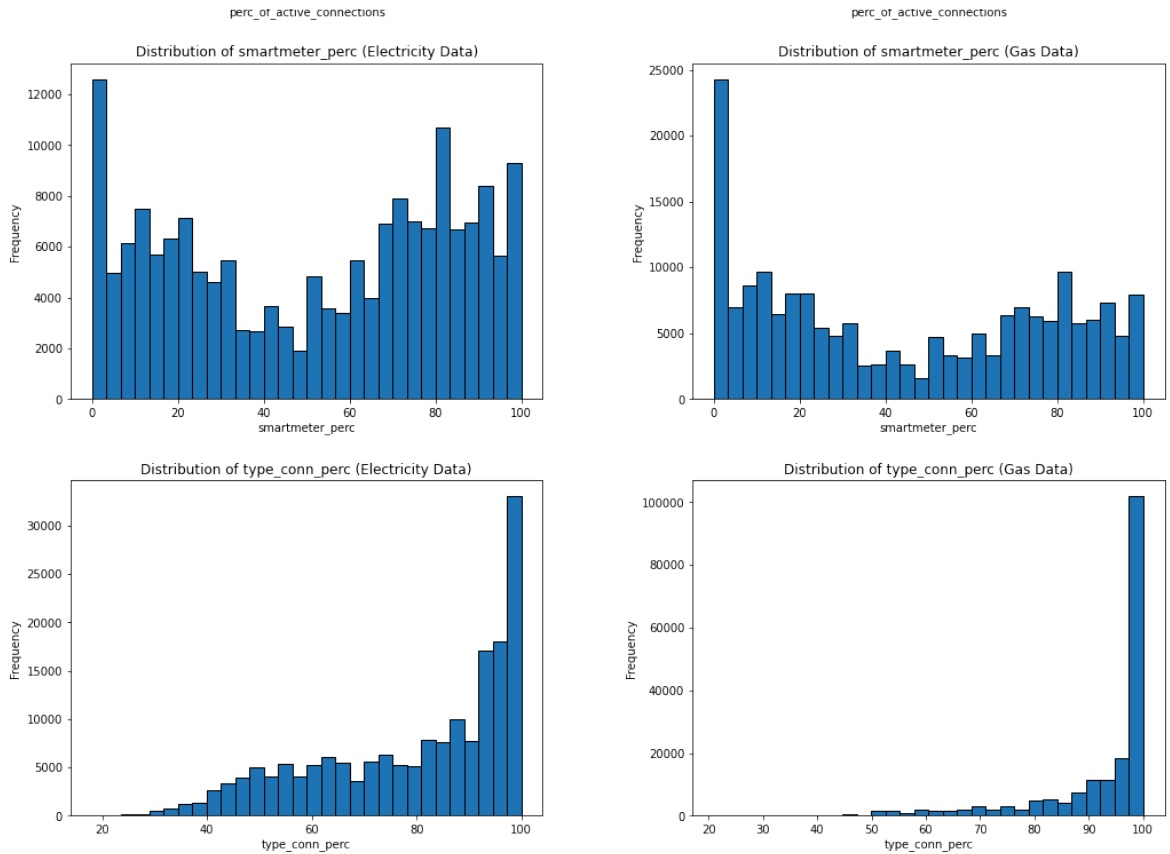
for i, col in enumerate(numerical_columns):
    # Electricity histogram
    axes[i, 0].hist(electricity_pandas_df[col], bins=30, edgecolor='black')
    axes[i, 0].set_title(f"Distribution of {col} (Electricity Data)")
    axes[i, 0].set_xlabel(col)
    axes[i, 0].set_ylabel("Frequency")

    # Gas histogram
    axes[i, 1].hist(gas_pandas_df[col], bins=30, edgecolor='black')
    axes[i, 1].set_title(f"Distribution of {col} (Gas Data)")
    axes[i, 1].set_xlabel(col)
    axes[i, 1].set_ylabel("Frequency")

plt.show()

```





Analyze the frequency distributions of key categorical columns to understand the diversity and dominance of specific categories. In the **city** column, both electricity and gas datasets have a significant number of unique values, indicating a high geographical coverage. The **net_manager** and **purchase_area** columns show skewed distributions, with a few dominant categories, such as **8716892000005** for **net_manager** and **Stedin** for **purchase_area**. Similarly, **type_of_connection** is dominated by **1x35** in electricity and **G4** in gas, indicating standardized types of connections.

```
In [0]: electricity_pandas_df["city"].value_counts()
```

```
Out[22]: 'S-GRAVENHAGE      23069
ROTTERDAM      20644
UTRECHT        11949
AMERSFOORT      5182
ZOETERMEER      5038
...
NIEUWGEIN       1
ZEVENHOVEN       1
'SâGRAVENHAGE    1
MAURIK           1
AMSTELVEEN       1
Name: city, Length: 287, dtype: int64
```

```
In [0]: gas_pandas_df["city"].value_counts()
```

```

Out[23]: 'S-GRAVENHAGE      20702
ROTTERDAM      17994
UTRECHT        9712
AMERSFOORT     4792
ZOETERMEER     4107
...
ZWAAGWESTEINDE      1
DAMWOUDE             1
NEDERHORST DEN BERG  1
'S GRAVENDEEL        1
NIGTEVEECHT         1
Name: city, Length: 491, dtype: int64

```

```
In [0]: electricity_pandas_df["net_manager"].value_counts()
```

```

Out[24]: 8716892000005      89266
8716874000009      51435
8716886000004      13073
8716921000006       8628
westland-infra      4842
8716925000002       3150
Cogas Infra & Beheer BV  2563
Coteq Netbeheer BV    2503
8716946000005       1331
8716924000003        16
Name: net_manager, dtype: int64

```

```
In [0]: gas_pandas_df["net_manager"].value_counts()
```

```

Out[25]: 8716892000005      81232
8716874000009      39433
8716886000004      11291
8716921000006       8015
Cogas Infra & Beheer BV  6868
Coteq Netbeheer BV     6774
8716892700004       6098
8716892750009       6032
8716892740000       5825
8716892710003       5134
westland-infra      4612
8716892720002       3285
8716925000002       3002
8716924000003         5
Name: net_manager, dtype: int64

```

```
In [0]: electricity_pandas_df["purchase_area"].value_counts()
```

```

Out[26]: Stedin      88881
Stedin Utrecht      51669
Stedin Delfland     13330
Stedin Midden-Holland  8524
Netbeheerder Centraal Overijssel B.V.  5066
871687800090000015    4842
Stedin Schiedam      3149
Stedin Elektriciteit Zuid-Kennemerland  1330
Stedin Weert         16
Name: purchase_area, dtype: int64

```

```
In [0]: gas_pandas_df["purchase_area"].value_counts()
```

```

Out[27]: NG Den Haag                26217
Pseudo Gos Houten ENBU             23313
Pseudo-GOS Dordrecht               17420
GAS Gastransport Services (GASUNIE) 13642
Pseudo-GOS Rotterdam               13100
Pseudo Gos Hoogland ENBU           9550
NG Leerdam                         8520
Pseudo-GOS Zoetermeer              6606
Pseudo Gos Veenendaal ENBU         6602
NG Gouda                           6313
Pseudo-GOS Zeist                   6032
Pseudo-GOS Amstelland              5580
NG Noord-Oost Friesland             5551
NG Hoekse waard                    5267
Pseudo-GOS Midden Kennemerland     5134
871718518003006694                 4612
Pseudo-GOS Delft                   3805
NG Heemstede                       3032
Schiedam Kethel                    3002
Pseudo-GOS Vlaardingen             2624
NG Brielle                         2313
NG Krimpen                         1664
NG Waddinxveen                     1662
Achterweg                          1610
Maassluis                          1077
Pseudo-GOS Bleiswijk               891
Oranjelaan                         460
Hoek van Holland                   386
Ouderkerk ad Amstel                366
Ameland                            272
Halfweg                            252
Ruigendijk                         228
Moerseweg                          217
Duivendrecht                       153
Graswalseweg                       115
Wildersekade                       12
Pseudo-GOS Weert                   4
NaN                                 1
Weert Trancheeweg                  1
Name: purchase_area, dtype: int64

```

```
In [0]: electricity_pandas_df["type_of_connection"].value_counts()
```

```

Out[28]: 1x35      126663
3x25      29887
1x25      14674
1x50       4416
3x35        769
3x63        189
3x80         98
3x50         85
OBK         19
1x6          7
Name: type_of_connection, dtype: int64

```


```
In [0]: gas_pandas_df["type_of_connection"].value_counts()
```

```
Out[29]: G4      182367
G6       4243
OBK      677
G16      216
G10       57
G25       46
Name: type_of_connection, dtype: int64
```


Generate correlation heatmaps for numerical columns to identify potential linear relationships between features. The result shows that most features have weak or no correlations.

```
In [0]: import seaborn as sns

# Heatmap for Electricity Data
plt.figure(figsize=(10, 8))
sns.heatmap(electricity_pandas_df[numerical_columns].corr(), annot=True, fmt=".2f",
plt.title("Correlation Heatmap (Electricity Data)")
plt.show()
```



```
In [0]: # Heatmap for Gas Data
plt.figure(figsize=(10, 8))
sns.heatmap(gas_pandas_df[numerical_columns].corr(), annot=True, fmt=".2f", cmap
plt.title("Correlation Heatmap (Gas Data)")
plt.show()
```



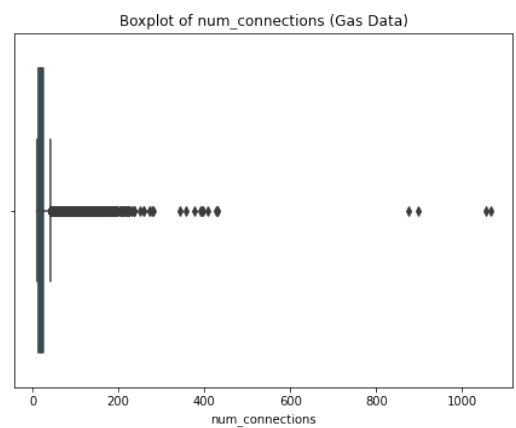
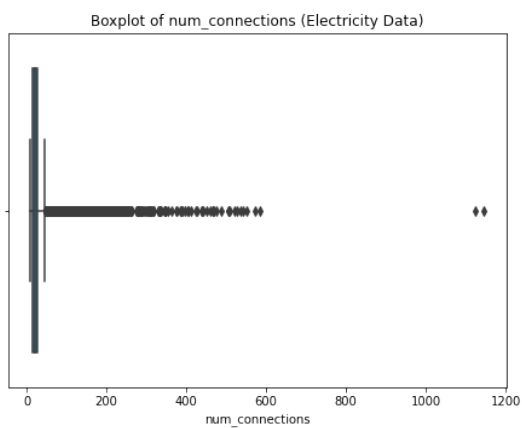
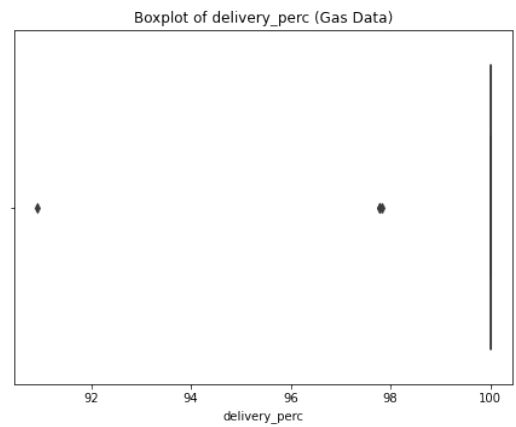
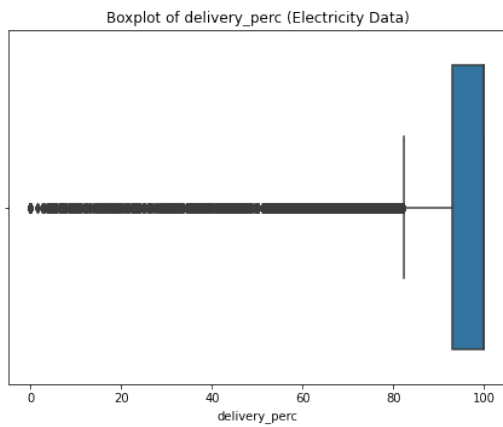
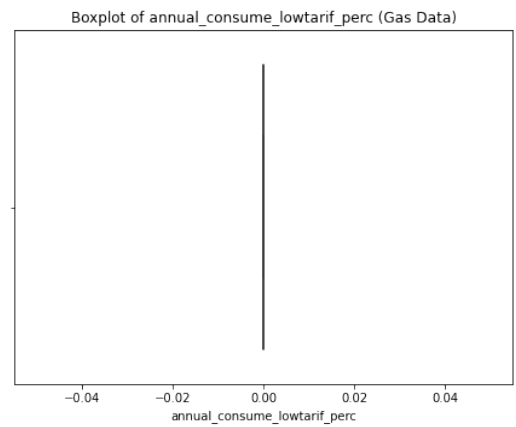
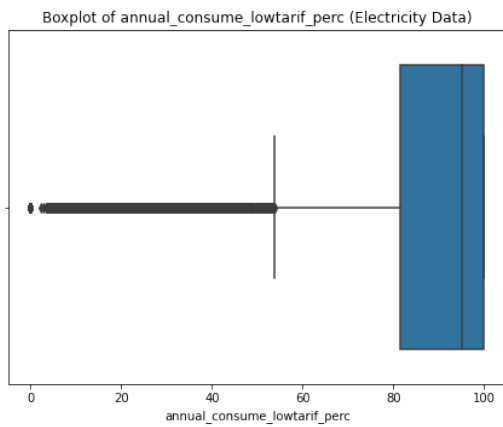
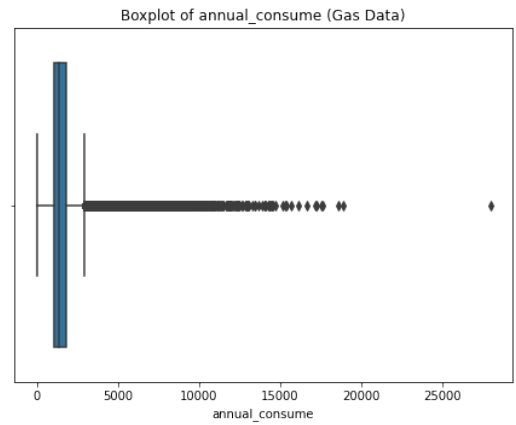
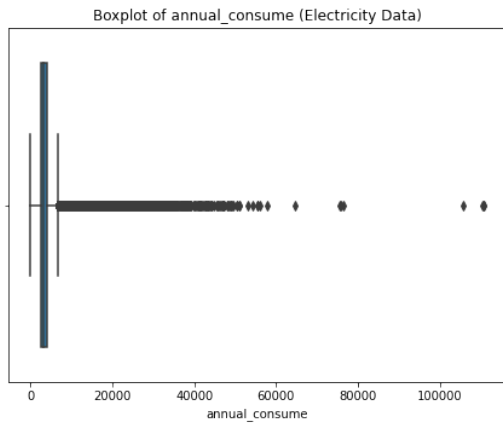
Generates boxplots for numerical columns to visualize the distributions and identify potential outliers. The boxplots reveal that several features, such as **annual_consume** and **num_connections**, have significant outliers in both datasets, indicating highly skewed distributions. For instance, **annual_consume** shows a long tail with extreme values, suggesting that these outliers may need to be handled during data cleaning to prevent them from negatively impacting model performance. Additionally, features like **delivery_perc** in the gas dataset appear to have no variability, as the values are constant, which might make the feature redundant for modeling and should be considered for removal.

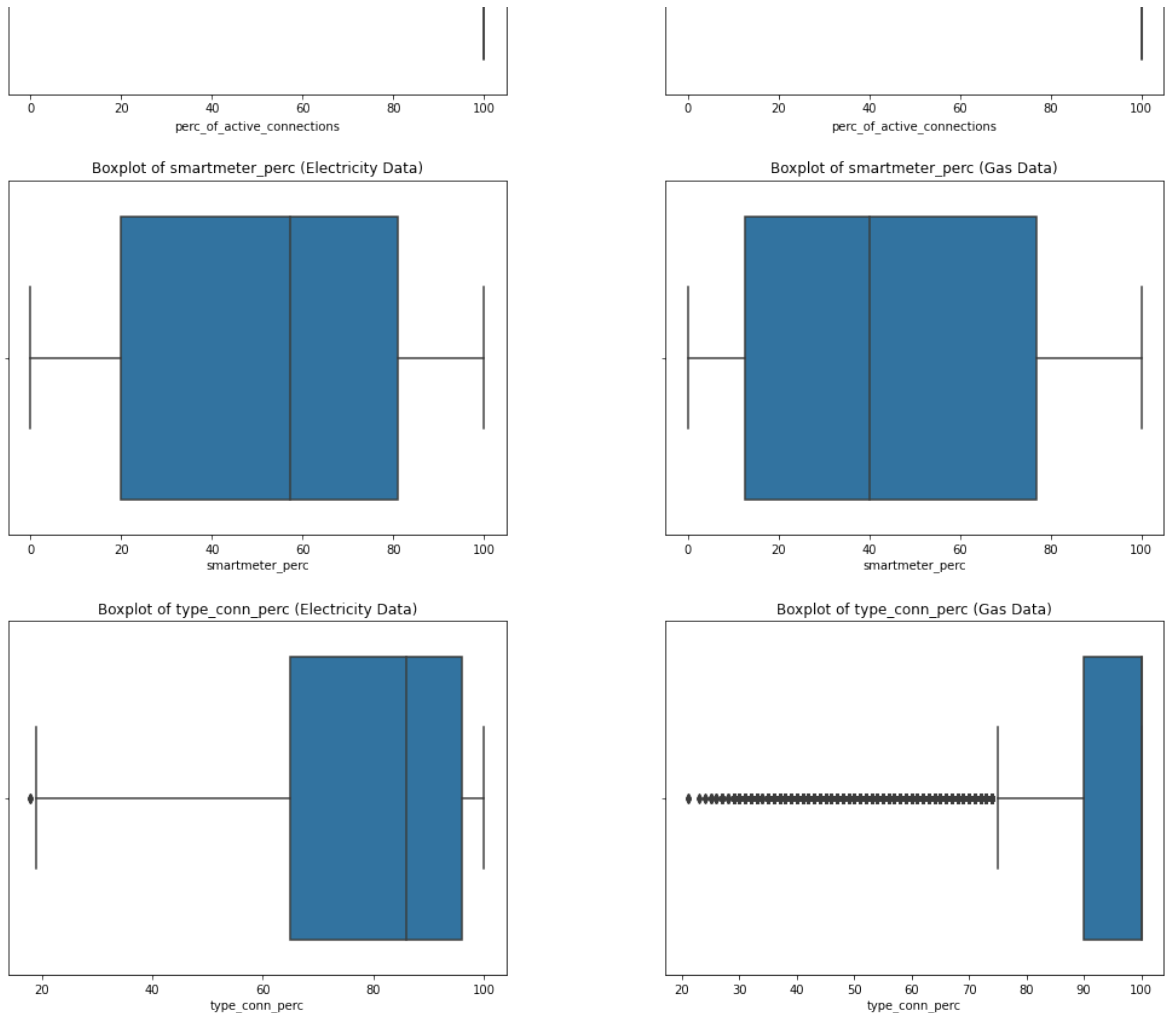
```
In [0]: fig, axes = plt.subplots(len(numerical_columns), 2, figsize=(14, 5 * len(numerical_columns)))
fig.tight_layout(w_pad=10.0, h_pad=5.0)

for i, col in enumerate(numerical_columns):
    sns.boxplot(x=electricity_pandas_df[col], ax=axes[i, 0])
    axes[i, 0].set_title(f"Boxplot of {col} (Electricity Data)")
    axes[i, 0].set_xlabel(col)

    sns.boxplot(x=gas_pandas_df[col], ax=axes[i, 1])
    axes[i, 1].set_title(f"Boxplot of {col} (Gas Data)")
    axes[i, 1].set_xlabel(col)

plt.show()
```





Question 2A: What are the chosen data cleaning steps? Why?

Step 1: Drop Columns with Missing Values

The **%Defintieve aansl (NRM)** column has an overwhelming number of missing values (171,965 in the electricity dataset and 182,994 in the gas dataset), accounting for the majority of rows. Imputation for such a large percentage of missing values would be impractical and unreliable, while retaining the column may introduce noise to the analysis. Since it does not provide meaningful information for modeling, this column is dropped entirely from all datasets.

```
In [0]: electricity_train_df = electricity_train_df.drop("%Defintieve aansl (NRM)")
gas_train_df = gas_train_df.drop("%Defintieve aansl (NRM)")
electricity_test_df = electricity_test_df.drop("%Defintieve aansl (NRM)")
gas_test_df = gas_test_df.drop("%Defintieve aansl (NRM)")
```

Step 2: Drop Columns with Low Analytical Values

Columns such as **_id**, **street**, **zipcode_from**, and **zipcode_to** are identifiers or high-cardinality features that do not provide meaningful patterns for analysis. Encoding them

would lead to sparse datasets, adding computational overhead without improving model accuracy. These columns are removed to focus on features with higher analytical value.

```
In [0]: columns_to_drop = ["_id", "street", "zipcode_from", "zipcode_to"]

electricity_train_df = electricity_train_df.drop(*columns_to_drop)
gas_train_df = gas_train_df.drop(*columns_to_drop)
electricity_test_df = electricity_test_df.drop(*columns_to_drop)
gas_test_df = gas_test_df.drop(*columns_to_drop)
```

Step 3: Drop Columns with Constant Values

The **annual_consume_lowtarif_perc** column in the gas dataset contains a constant value of 0 across all rows, offering no variability or predictive power for the model. Retaining this column would introduce redundancy without contributing to the model's performance. Thus, this column is removed to simplify the dataset and improve efficiency.

```
In [0]: gas_train_df = gas_train_df.drop("annual_consume_lowtarif_perc")
gas_test_df = gas_test_df.drop("annual_consume_lowtarif_perc")
```

Step 4: Handle Outliers in Numerical Columns

Exploratory analysis revealed extreme outliers in numerical columns such as **annual_consume**, **num_connections**, and **type_conn_perc**, with maximum values significantly exceeding the mean. These outliers distort statistical summaries and can negatively affect model training. To address this, log transformation is applied to these columns, reducing skewness and stabilizing variance, ensuring the data is more suitable for predictive modeling.

```
In [0]: from pyspark.sql.functions import log1p

columns_to_transform = ["annual_consume", "num_connections", "type_conn_perc"]
for col in columns_to_transform:
    electricity_train_df = electricity_train_df.withColumn(col, log1p(electricit
    gas_train_df = gas_train_df.withColumn(col, log1p(gas_train_df[col]))
    electricity_test_df = electricity_test_df.withColumn(col, log1p(electricity_
    gas_test_df = gas_test_df.withColumn(col, log1p(gas_test_df[col]))
```

Step 5: Verify Data Integrity

After cleaning, it is essential to ensure data integrity by checking for duplicate rows, missing values, or anomalies in the datasets. Duplicate rows are removed to prevent biases during modeling, and missing value checks confirm that no additional imputation or removal is required. This step ensures the datasets are fully prepared for transformation and analysis.

```
In [0]: # Remove duplicates
electricity_train_df = electricity_train_df.dropDuplicates()
gas_train_df = gas_train_df.dropDuplicates()
```

```
electricity_test_df = electricity_test_df.dropDuplicates()
gas_test_df = gas_test_df.dropDuplicates()
```

```
In [0]: from pyspark.sql.functions import col, sum as _sum

# Verify missing values
display(electricity_train_df.select([_sum(col(c).isNull().cast("int")).alias(c)
```

annual_consume	annual_consume_lowtarif_perc	city	delivery_perc	net_manager	num
0	0	0	0	0	0

```
In [0]: display(gas_train_df.select([_sum(col(c).isNull().cast("int")).alias(c) for c in
```

annual_consume	city	delivery_perc	net_manager	num_connections	perc_of_active_co
0	0	0	0	0	0

Question 2B: What are the chosen data transformation steps? Why?

Step 1: Encode Categorical Variables

Categorical variables, such as **city**, **net_manager**, **purchase_area**, and **type_of_connection**, are non-numeric and cannot be directly used by most machine learning algorithms. Using `StringIndexer`, each category is assigned a unique numerical index, capturing categorical distinctions. To avoid introducing ordinal bias (where numerical values imply ranking), these indices are converted into binary vectors using `OneHotEncoder`. This step ensures that all categorical variables are represented in a way that preserves their relationships without misinterpreting their values as numeric magnitudes.

```
In [0]: from pyspark.ml.feature import StringIndexer, OneHotEncoder

categorical_columns = ["city", "net_manager", "purchase_area", "type_of_connecti

# String Indexing and One-Hot Encoding for each categorical column
indexers = [StringIndexer(inputCol=col, outputCol=f"{col}_index", handleInvalid=
encoders = [OneHotEncoder(inputCol=f"{col}_index", outputCol=f"{col}_onehot") fo

# Combine all transformations into a List
categorical_transformations = indexers + encoders
```

Step 2: Scale Numerical Features

Numerical features, such as **num_connections** and **type_conn_perc**, have varying scales and units, as observed during data exploration. Features with larger magnitudes may disproportionately influence model training, especially for algorithms like decision trees.

To mitigate this, `MinMaxScaler` is applied to normalize all numerical features to a range of [0, 1], ensuring equal contribution of all variables during training.

```
In [0]: from pyspark.ml.feature import MinMaxScaler, VectorAssembler

numerical_columns = ["num_connections", "type_conn_perc"]

# Assemble numerical columns into a single vector for scaling
assembler = VectorAssembler(inputCols=numerical_columns, outputCol="numerical_features")
scaler = MinMaxScaler(inputCol="numerical_features", outputCol="scaled_numerical_features")

# Combine assembler and scaler into a list
numerical_transformations = [assembler, scaler]
```

Step 3: Assemble Features

Machine learning models typically require a single vector column containing all the features. After encoding categorical variables and scaling numerical features, the `VectorAssembler` consolidates all transformed columns (scaled_numerical_features and one-hot encoded categorical columns) into a unified features column. This step ensures that the data is structured consistently and efficiently for input into machine learning pipelines.

```
In [0]: all_features = [f"{col}_onehot" for col in categorical_columns] + ["scaled_numerical_features"]

# Assemble all features into a single vector
feature_assembler = VectorAssembler(inputCols=all_features, outputCol="features")
```

Task 3: Model training and tracking with data pipeline and MLflow

Data Preparation

```
In [0]: # Target column
target_column = "annual_consume"

# Prepare datasets with labeled target column
electricity_train_df = electricity_train_df.withColumnRenamed(target_column, "label")
electricity_test_df = electricity_test_df.withColumnRenamed(target_column, "label")
gas_train_df = gas_train_df.withColumnRenamed(target_column, "label")
gas_test_df = gas_test_df.withColumnRenamed(target_column, "label")
```

Define Reusable Training and Logging Pipeline

```
In [0]: from pyspark.ml import Pipeline
from pyspark.ml.regression import RandomForestRegressor, DecisionTreeRegressor,
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.sql.functions import col
import mlflow
import mlflow.spark
```

```

def train_and_log_model(train_data, test_data, algorithm, params):
    with mlflow.start_run():
        # Initialize the model
        if algorithm == "RandomForestRegressor":
            model = RandomForestRegressor(featuresCol="features", labelCol="label")
        elif algorithm == "DecisionTreeRegressor":
            model = DecisionTreeRegressor(featuresCol="features", labelCol="label")

        # Create pipeline
        pipeline = Pipeline(stages=categorical_transformations + numerical_transformations)

        # Train the model
        trained_pipeline = pipeline.fit(train_data)

        # Predict on test data
        predictions = trained_pipeline.transform(test_data)

        # Evaluate the model
        evaluator_mae = RegressionEvaluator(labelCol="label", predictionCol="prediction", metricName="mae")
        evaluator_r2 = RegressionEvaluator(labelCol="label", predictionCol="prediction", metricName="r2")
        evaluator_rmse = RegressionEvaluator(labelCol="label", predictionCol="prediction", metricName="rmse")

        mae = evaluator_mae.evaluate(predictions)
        r2 = evaluator_r2.evaluate(predictions)
        rmse = evaluator_rmse.evaluate(predictions)

        # Log parameters, metrics, and the model
        mlflow.log_param("algorithm", algorithm)
        mlflow.log_params(params)
        mlflow.log_metric("MAE", mae)
        mlflow.log_metric("R2", r2)
        mlflow.log_metric("RMSE", rmse)
        mlflow.spark.log_model(trained_pipeline, "model")

        print(f"Logged {algorithm} with params {params}")
        print("Evaluation metrics:")
        print(f"- MAE: {mae}")
        print(f"- R2: {r2}")
        print(f"- RMSE: {rmse}")

```

Define Different Parameter Settings for Each Algorithm.

```

In [0]: # Random Forest parameters
random_forest_params = [
    {"numTrees": 10, "maxDepth": 3}, # Small number of trees, shallow depth
    {"numTrees": 20, "maxDepth": 5}, # Slightly deeper with more trees
    {"numTrees": 30, "maxDepth": 7} # Moderate depth and tree count
]

# Decision Tree parameters
decision_tree_params = [
    {"maxDepth": 3, "minInstancesPerNode": 1}, # Shallow tree
    {"maxDepth": 5, "minInstancesPerNode": 2}, # Slightly deeper
    {"maxDepth": 7, "minInstancesPerNode": 2} # Moderate depth
]

```

Electricity Model Training

Random Forest Regressor Model

```
In [0]: train_and_log_model(electricity_train_df, electricity_test_df, "RandomForestRegrn
```

```
Logged RandomForestRegressor with params {'numTrees': 10, 'maxDepth': 3}
Evaluation metrics:
- MAE: 0.3245205459858661
- R2: 0.3055551751803406
- RMSE: 0.43719801389464497
```

```
In [0]: train_and_log_model(electricity_train_df, electricity_test_df, "RandomForestRegrn
```

```
Logged RandomForestRegressor with params {'numTrees': 20, 'maxDepth': 5}
Evaluation metrics:
- MAE: 0.30309765928996096
- R2: 0.4100154729258382
- RMSE: 0.4029764052916728
```

```
In [0]: train_and_log_model(electricity_train_df, electricity_test_df, "RandomForestRegrn
```

```
Logged RandomForestRegressor with params {'numTrees': 30, 'maxDepth': 7}
Evaluation metrics:
- MAE: 0.2924566271932784
- R2: 0.45467513836524165
- RMSE: 0.38742438864546014
```

Decision Tree Regressor Model

```
In [0]: train_and_log_model(electricity_train_df, electricity_test_df, "DecisionTreeRegrn
```

```
Logged DecisionTreeRegressor with params {'maxDepth': 3, 'minInstancesPerNode': 1}
Evaluation metrics:
- MAE: 0.3198013444238246
- R2: 0.33819237191648355
- RMSE: 0.4268007675394118
```

```
In [0]: train_and_log_model(electricity_train_df, electricity_test_df, "DecisionTreeRegrn
```

```
Logged DecisionTreeRegressor with params {'maxDepth': 5, 'minInstancesPerNode': 2}
Evaluation metrics:
- MAE: 0.3005998752347698
- R2: 0.41813901313332635
- RMSE: 0.40019248323442663
```

```
In [0]: train_and_log_model(electricity_train_df, electricity_test_df, "DecisionTreeRegrn
```

```
Logged DecisionTreeRegressor with params {'maxDepth': 7, 'minInstancesPerNode': 2}
Evaluation metrics:
- MAE: 0.2919387583533915
- R2: 0.4505358781081866
- RMSE: 0.3888919714346681
```


Gas Model Training

Random Forest Regressor Model

```
In [0]: train_and_log_model(gas_train_df, gas_test_df, "RandomForestRegressor", random_f
```

```
Logged RandomForestRegressor with params {'numTrees': 10, 'maxDepth': 3}
Evaluation metrics:
- MAE: 0.4311770564844025
- R2: 0.1673514212579711
- RMSE: 0.7598437611829207
```

```
In [0]: train_and_log_model(gas_train_df, gas_test_df, "RandomForestRegressor", random_f
```

```
Logged RandomForestRegressor with params {'numTrees': 20, 'maxDepth': 5}
Evaluation metrics:
- MAE: 0.4384877355674457
- R2: 0.06385954112566983
- RMSE: 0.8056825121370964
```

```
In [0]: train_and_log_model(gas_train_df, gas_test_df, "RandomForestRegressor", random_f
```

```
Logged RandomForestRegressor with params {'numTrees': 30, 'maxDepth': 7}
Evaluation metrics:
- MAE: 0.5128250100815277
- R2: 0.12755182278567023
- RMSE: 0.7777916012613559
```

Decision Tree Regressor Model

```
In [0]: train_and_log_model(gas_train_df, gas_test_df, "DecisionTreeRegressor", decision
```

```
Logged DecisionTreeRegressor with params {'maxDepth': 3, 'minInstancesPerNode':
1}
Evaluation metrics:
- MAE: 0.4292077809910484
- R2: 0.17396052800107087
- RMSE: 0.7568221421899121
```

```
In [0]: train_and_log_model(gas_train_df, gas_test_df, "DecisionTreeRegressor", decision
```

```
Logged DecisionTreeRegressor with params {'maxDepth': 5, 'minInstancesPerNode':
2}
Evaluation metrics:
- MAE: 0.4210915936830995
- R2: 0.19771574365696465
- RMSE: 0.745860425127864
```

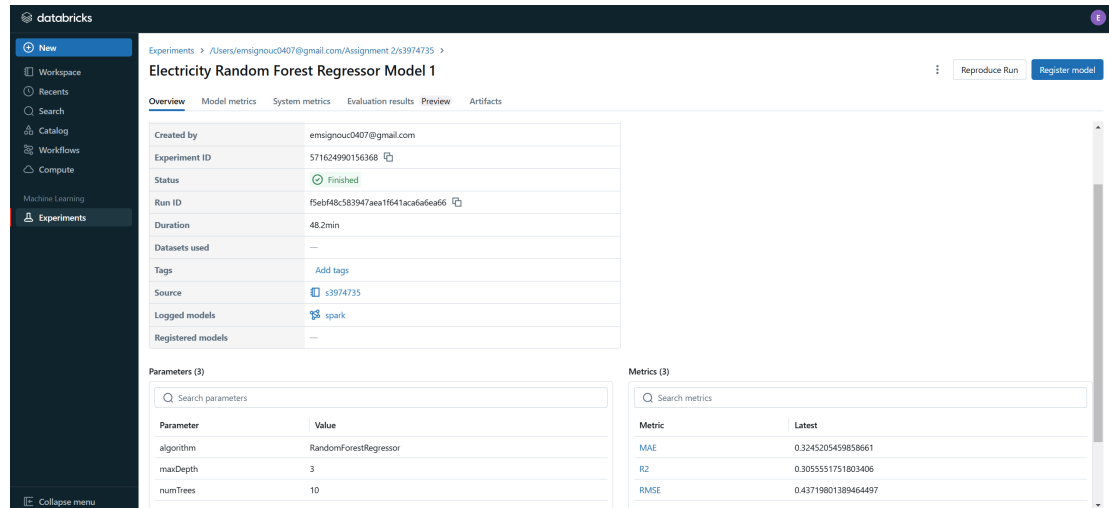
```
In [0]: train_and_log_model(gas_train_df, gas_test_df, "DecisionTreeRegressor", decision
```

```
Logged DecisionTreeRegressor with params {'maxDepth': 7, 'minInstancesPerNode':
2}
Evaluation metrics:
- MAE: 0.4177651838941951
- R2: 0.20934962369445542
- RMSE: 0.7404328364065984
```

Question 3A: What is/are your final model(s) based on the evaluation metrics?

Electricity Model

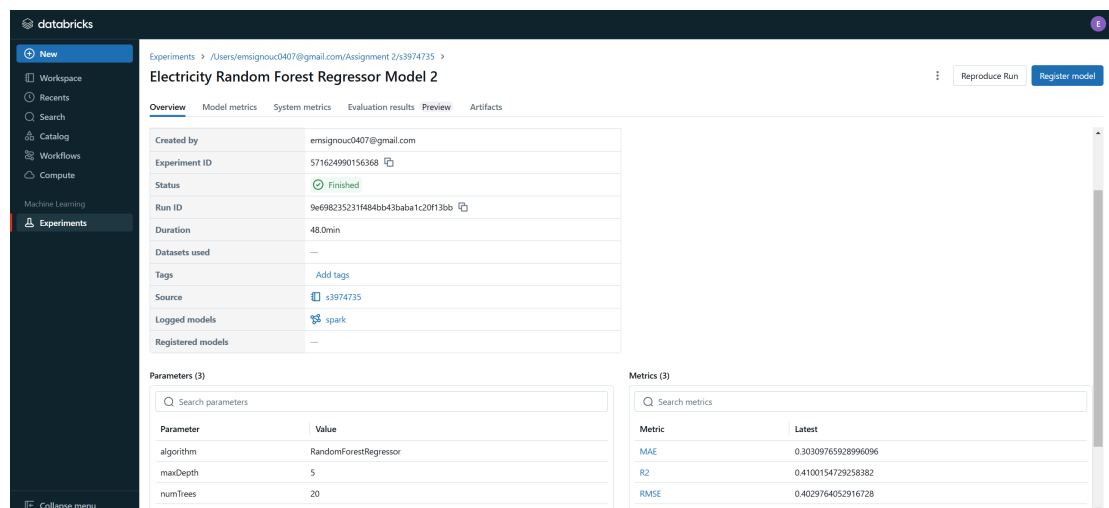
MLflow UI for Electricity Random Forest Regressor Models:



This screenshot shows the MLflow UI for 'Electricity Random Forest Regressor Model 1'. The interface includes a sidebar with navigation options like 'New', 'Workspace', 'Recents', 'Search', 'Catalog', 'Workflows', 'Compute', 'Machine Learning', and 'Experiments'. The main content area displays the experiment details for 'Electricity Random Forest Regressor Model 1' (Experiment ID: 571624990156368). The 'Overview' tab is selected, showing a table of parameters and a table of metrics.

Parameter	Value
algorithm	RandomForestRegressor
maxDepth	3
numTrees	10

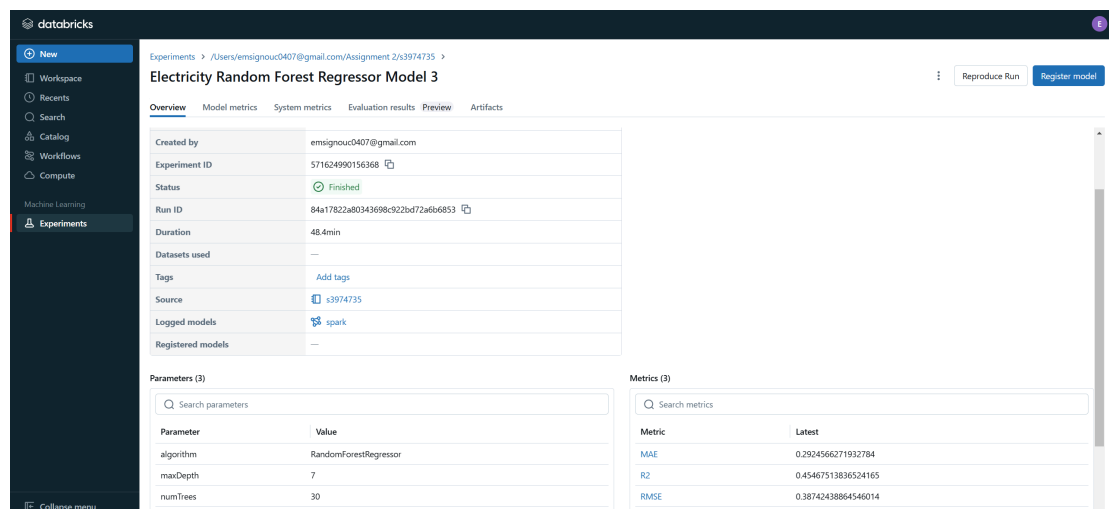
Metric	Latest
MAE	0.3245205459858661
R2	0.3055551751803406
RMSE	0.4371980138946497



This screenshot shows the MLflow UI for 'Electricity Random Forest Regressor Model 2' (Experiment ID: 571624990156368). The 'Overview' tab displays the experiment details, including parameters and metrics.

Parameter	Value
algorithm	RandomForestRegressor
maxDepth	5
numTrees	20

Metric	Latest
MAE	0.30309765928996096
R2	0.4100154729258382
RMSE	0.4029764052916728

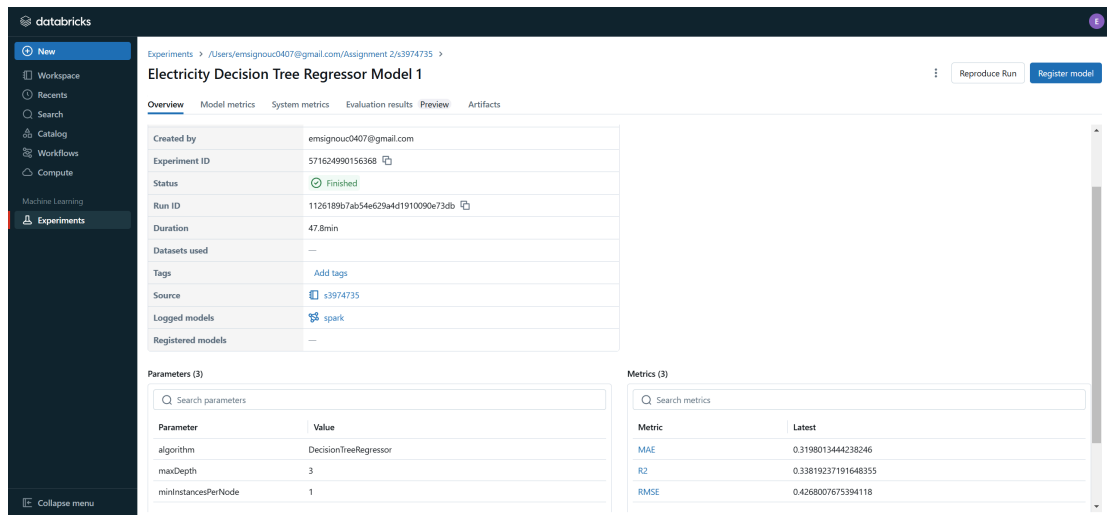


This screenshot shows the MLflow UI for 'Electricity Random Forest Regressor Model 3' (Experiment ID: 571624990156368). The 'Overview' tab displays the experiment details, including parameters and metrics.

Parameter	Value
algorithm	RandomForestRegressor
maxDepth	7
numTrees	30

Metric	Latest
MAE	0.2924566271932784
R2	0.45467513836524165
RMSE	0.38742438864546014

MLflow UI for Electricity Decision Tree Regressor Models:



Electricity Decision Tree Regressor Model 1

Overview | Model metrics | System metrics | Evaluation results | Preview | Artifacts

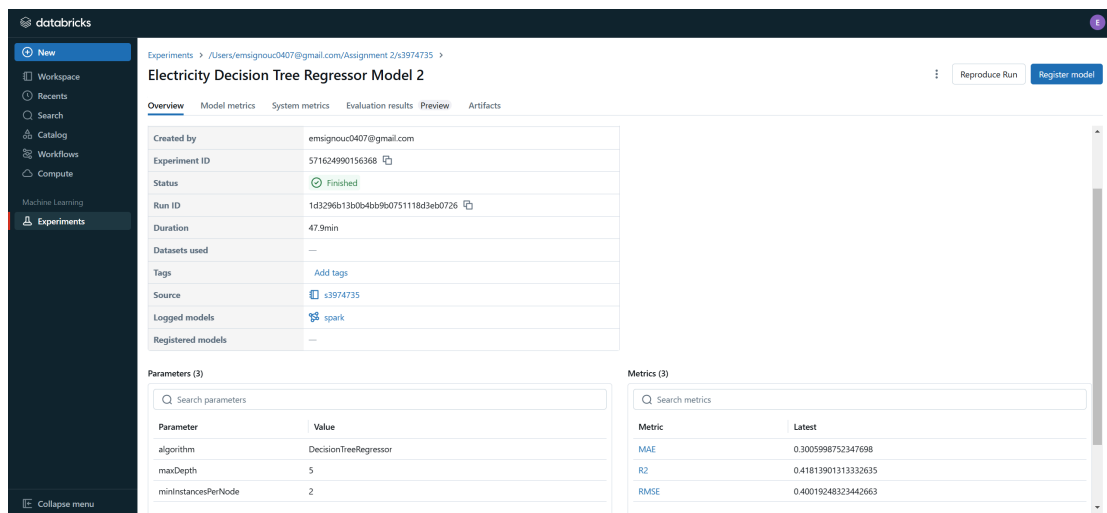
Created by: emsignou0407@gmail.com
Experiment ID: 571624990156368
Status: Finished
Run ID: 1126189b7ab54e629a4d1910090e73db
Duration: 47.8min
Datasets used: —
Tags: [Add tags](#)
Source: [s3974735](#)
Logged models: [spark](#)
Registered models: —

Parameters (3)

Parameter	Value
algorithm	DecisionTreeRegressor
maxDepth	3
minInstancesPerNode	1

Metrics (3)

Metric	Latest
MAE	0.3198013444238246
R2	0.33819237191646355
RMSE	0.4268007675394118



Electricity Decision Tree Regressor Model 2

Overview | Model metrics | System metrics | Evaluation results | Preview | Artifacts

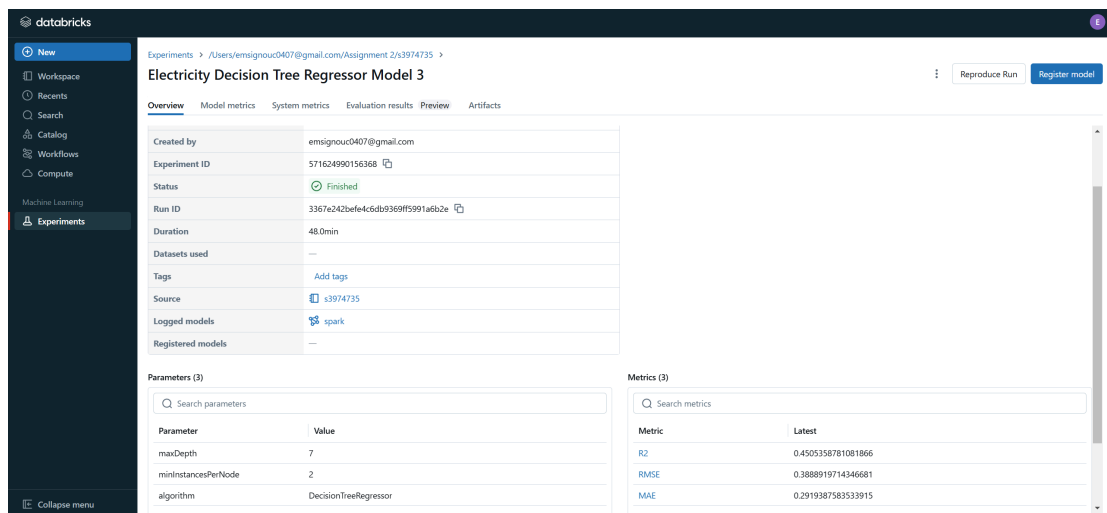
Created by: emsignou0407@gmail.com
Experiment ID: 571624990156368
Status: Finished
Run ID: 1d3296b13b0b4bb960751118d3eb0726
Duration: 47.9min
Datasets used: —
Tags: [Add tags](#)
Source: [s3974735](#)
Logged models: [spark](#)
Registered models: —

Parameters (3)

Parameter	Value
algorithm	DecisionTreeRegressor
maxDepth	5
minInstancesPerNode	2

Metrics (3)

Metric	Latest
MAE	0.3005998752347698
R2	0.4181390131332635
RMSE	0.40019248323442663



Electricity Decision Tree Regressor Model 3

Overview | Model metrics | System metrics | Evaluation results | Preview | Artifacts

Created by: emsignou0407@gmail.com
Experiment ID: 571624990156368
Status: Finished
Run ID: 3367e242befe4c5db9369ff5901a6b2e
Duration: 48.0min
Datasets used: —
Tags: [Add tags](#)
Source: [s3974735](#)
Logged models: [spark](#)
Registered models: —

Parameters (3)

Parameter	Value
maxDepth	7
minInstancesPerNode	2
algorithm	DecisionTreeRegressor

Metrics (3)

Metric	Latest
R2	0.4505358781081866
RMSE	0.3888919714346681
MAE	0.2919387583533915

Based on the evaluation metrics, the best-performing model for the electricity dataset is the **RandomForestRegressor** with parameters `{ 'numTrees': 30, 'maxDepth': 7 }`. This model achieved the lowest MAE of 0.2925, the highest R2 of 0.4547, and the lowest RMSE of 0.3874 among all tested models and parameter settings. These metrics indicate that this configuration provides the most accurate predictions with the least error and the best explained variance.

The **DecisionTreeRegressor** with parameters `{ 'maxDepth': 7, 'minInstancesPerNode': 2 }` performed closely, achieving an MAE of 0.2919, an R2 of

0.4505, and an RMSE of 0.3889. However, the RandomForestRegressor slightly outperformed it, likely due to its ensemble nature, which reduces variance and enhances prediction stability.

Thus, the final model for the electricity dataset is the **RandomForestRegressor** with `{ 'numTrees': 30, 'maxDepth': 7 }` due to its superior performance across all evaluation metrics.

Gas Model

MLflow UI for Gas Random Forest Regressor Models:

New

Workspace

Recents

Search

Catalog

Workflows

Compute

Machine Learning

Experiments

Experiments

/Users/emsignou0407@gmail.com/Assignment 2/s3974735

Gas Random Forest Regressor Model 1

Reproduce Run

Register model

Overview

Model metrics

System metrics

Evaluation results

Preview

Artifacts

Created by

emsignou0407@gmail.com

Experiment ID

571624990156368

Status

Finished

Run ID

acd8809e7c184ac5ae7986ad5ba9aef2

Duration

45.1min

Datasets used

—

Tags

Add tags

Source

s3974735

Logged models

spark

Registered models

—

Parameters (3)

Search parameters

Parameter

Value

algorithm

RandomForestRegressor

maxDepth

3

numTrees

10

Metrics (3)

Search metrics

Metric

Latest

MAE

0.4311770564844025

R2

0.1673514212579711

RMSE

0.7598437611829207

New

Workspace

Recents

Search

Catalog

Workflows

Compute

Machine Learning

Experiments

Experiments

/Users/emsignou0407@gmail.com/Assignment 2/s3974735

Gas Random Forest Regressor Model 2

Reproduce Run

Register model

Overview

Model metrics

System metrics

Evaluation results

Preview

Artifacts

Created by

emsignou0407@gmail.com

Experiment ID

571624990156368

Status

Finished

Run ID

cbb86e2b1ee8444763eeb1645a8bdad0

Duration

45.3min

Datasets used

—

Tags

Add tags

Source

s3974735

Logged models

spark

Registered models

—

Parameters (3)

Search parameters

Parameter

Value

numTrees

20

maxDepth

5

algorithm

RandomForestRegressor

Metrics (3)

Search metrics

Metric

Latest

RMSE

0.8056825121370964

R2

0.06385954112566983

MAE

0.4384877355674457

databricks

New

Workspace

Recents

Search

Catalog

Workflows

Compute

Machine Learning

Experiments

Collapse menu

Experiments > /Users/emsignouc0407@gmail.com/Assignment 2/s3974735 >

Gas Random Forest Regressor Model 3

Reproduce Run

Register model

Overview

Model metrics

System metrics

Evaluation results

Preview

Artifacts

Created by

emsignouc0407@gmail.com

Experiment ID

571624990156368

Status

Finished

Run ID

f147ae5a8eb458cb0ad300ecfa6ab8

Duration

45.7min

Datasets used

—

Tags

Add tags

Source

s3974735

Logged models

spark

Registered models

—

Parameters (3)

Search parameters

Parameter	Value
numTrees	30
maxDepth	7
algorithm	RandomForestRegressor

Metrics (3)

Search metrics

Metric	Latest
RMSE	0.7777916012613559
R2	0.12755182278567023
MAE	0.5128250100815277

MLflow UI for Gas Decision Tree Regressor Models:

databricks

New

Workspace

Recents

Search

Catalog

Workflows

Compute

Machine Learning

Experiments

Collapse menu

Experiments > /Users/emsignouc0407@gmail.com/Assignment 2/s3974735 >

Gas Decision Tree Regressor Model 1

Reproduce Run

Register model

Overview

Model metrics

System metrics

Evaluation results

Preview

Artifacts

Created by

emsignouc0407@gmail.com

Experiment ID

571624990156368

Status

Finished

Run ID

d93257cc58364078a666b436177ee197

Duration

44.9min

Datasets used

—

Tags

Add tags

Source

s3974735

Logged models

spark

Registered models

—

Parameters (3)

Search parameters

Parameter	Value
minInstancesPerNode	1
maxDepth	3
algorithm	DecisionTreeRegressor

Metrics (3)

Search metrics

Metric	Latest
RMSE	0.7568221421899121
R2	0.17396052800107087
MAE	0.4292077809910484

databricks

New

Workspace

Recents

Search

Catalog

Workflows

Compute

Machine Learning

Experiments

Collapse menu

Experiments > /Users/emsignouc0407@gmail.com/Assignment 2/s3974735 >

Gas Decision Tree Regressor Model 2

Reproduce Run

Register model

Overview

Model metrics

System metrics

Evaluation results

Preview

Artifacts

Created by

emsignouc0407@gmail.com

Experiment ID

571624990156368

Status

Finished

Run ID

3e58e8215f12422e9d55baafcc02ad5f

Duration

44.9min

Datasets used

—

Tags

Add tags

Source

s3974735

Logged models

spark

Registered models

—

Parameters (3)

Search parameters

Parameter	Value
algorithm	DecisionTreeRegressor
maxDepth	5
minInstancesPerNode	2

Metrics (3)

Search metrics

Metric	Latest
MAE	0.4210915936830995
R2	0.19771574365696465
RMSE	0.745860425127864

file:///D:/Projects/mongodb_and_spark/s3974735.html

29/33

Gas Decision Tree Regressor Model 3

Overview | Model metrics | System metrics | Evaluation results | Preview | Artifacts

Created by	emsignou0407@gmail.com
Experiment ID	571624990156368
Status	Finished
Run ID	fbcb5cde8b14839810515ae8042f888
Duration	44.9min
Datasets used	—
Tags	Add tags
Source	s3974735
Logged models	spark
Registered models	—

Parameters (3)

Parameter	Value
minInstancesPerNode	2
maxDepth	7
algorithm	DecisionTreeRegressor

Metrics (3)

Metric	Latest
RMSE	0.7404328364065984
R2	0.20934962369445542
MAE	0.4177651838941951

Based on the evaluation metrics, the best-performing model for the gas dataset is the **DecisionTreeRegressor** with parameters `{ 'maxDepth': 7, 'minInstancesPerNode': 2 }`. This model achieved the lowest MAE of 0.4178, the highest R2 of 0.2093, and the lowest RMSE of 0.7404 among all tested configurations. These metrics suggest that this model has the best balance of accuracy, explained variance, and minimal error compared to other tested models.

Although the **RandomForestRegressor** was evaluated with multiple configurations, its performance was consistently inferior to the **DecisionTreeRegressor**. The highest R2 achieved by the **RandomForestRegressor** was only 0.1674, with an RMSE of 0.7598 and an MAE of 0.4312, indicating lower predictive accuracy and higher error rates.

Thus, the final model for the gas dataset is the **DecisionTreeRegressor** with parameters `{ 'maxDepth': 7, 'minInstancesPerNode': 2 }` due to its superior performance across all evaluation metrics.

Question 3B: Did you build one model for both electricity and gas or separate models? Why?

Separate models were built for the electricity and gas datasets. This decision stems from the insights gained during data exploration, which revealed significant differences between the two datasets:

- The **annual_consume** column in the electricity dataset has a much higher mean of 3913 and maximum value of 110,857 compared to the gas dataset, where the mean is 1559 and the maximum is 27,917. This indicates fundamentally different consumption behaviors between the two datasets.
- Numerical columns such as **num_connections** and **type_conn_perc** exhibit distinct patterns in the two datasets. Electricity data shows higher variability in **num_connections**, with values reaching 1146, while gas data has a maximum of 1065. Similarly, **type_conn_perc** is generally higher in the gas dataset.
- Categorical columns like **city**, **net_manager**, and **purchase_area** have varying levels of diversity and dominant categories. For instance, the electricity dataset has 287

unique cities, while the gas dataset has 491. Additionally, the **purchase_area** column features entirely different dominant categories for each dataset.

- Correlation heatmaps showed weak or no relationships between most numerical features in both datasets. The lack of shared strong correlations indicates that the datasets would likely require different feature importance considerations during model training.
- The **annual_consume_lowtarif_perc** column in the gas dataset is constant with a value of 0.0 and adds no value to model predictions, whereas it has variability in the electricity dataset.

These differences demonstrate that the datasets represent distinct consumption patterns, influenced by their respective geographical, infrastructural, and operational contexts. Building separate models ensures that each model is tailored to the unique characteristics of its dataset, thereby improving the accuracy, robustness, and interpretability of the predictions.

Question 3C: Should we build a separate model for each company or not? Why?

Although for this project we are working with the datasets of three different companies, it is unnecessary to create separate models for each company. This decision is supported by several observations:

- During data exploration, it was found that while the **net_manager** and **purchase_area** columns influence data distribution, they do not exhibit substantial variability in their effect on the target variable, **annual_consume**. The target variable is more strongly influenced by features like **num_connections**, **delivery_perc**, and **type_conn_perc**, which are not tied to specific companies.
- **net_manager** and **purchase_area** columns also have high cardinality and sparse representation for some companies. Building separate models would result in less-represented companies having sparse datasets, potentially causing overfitting and reducing the generalizability of the models.
- Key features such as **annual_consume_lowtarif_perc** and **smartmeter_perc** display similar predictive relationships across companies. This indicates that a single model can effectively generalize across companies without requiring specific adjustments.
- Creating separate models for each company would increase computational complexity, development time, and maintenance effort. A single model for the electricity dataset and another for the gas dataset ensures a scalable and efficient solution while maintaining strong predictive accuracy.
- By including company-specific features like **net_manager** as categorical inputs in the models, the differences between companies can still be accounted for without needing separate models.

For these reasons, building one model for each dataset is sufficient to achieve accurate predictions while keeping the solution simple, scalable, and easy to manage.

Task 4: Visualisation

Link to MongoDB Charts Dashboard: <https://charts.mongodb.com/charts-bigdataasm2-szigrao/public/dashboards/67724c59-0c78-4054-8e6b-1061df46332b>

Chart 1: Top 10 Cities by Electricity Annual Consumption in 2018




Chart 2: Electricity Distribution of Connection Types in 2019




Chart 3: Top 10 Cities by Gas Annual Consumption in 2018




Chart 4: Gas Distribution of Connection Types in 2019

